

Almost-Linear Time Algorithms for Decremental Graphs: Min-Cost Flow and More via Duality

Li Chen

Stanford 05/17/24

Joint work with



Jan van den
Brand
Georgia Tech



Rasmus Kyng
ETH



Yang P. Liu
IAS



Simon
Meierhans
ETH




Maximilian
Probst Gutenberg
ETH



Sushant
Sachdeva
U. Toronto

Uncapacitated Min-Cost Flow


Directed graph $G = (V, E)$. m edges, n vertices, source s , sink t
edge costs c_e , integer in $[-C, C]$, where $C = m^{O(1)}$



Uncapacitated Min-Cost Flow

Directed graph $G = (V, E)$. m edges, n vertices, source s , sink t
edge costs c_e , integer in $[-C, C]$, where $C = m^{O(1)}$

Goal: Route 1 unit of
flow from $s \rightarrow t$,
minimize cost




Uncapacitated Min-Cost Flow

Directed graph $G = (V, E)$. m edges, n vertices, source s , sink t
edge costs c_e , integer in $[-C, C]$, where $C = m^{O(1)}$

Goal: Route 1 unit of
flow from $s \rightarrow t$,
minimize cost

Real-weighted shortest
st-path




Uncapacitated Min-Cost Flow

Directed graph $G = (V, E)$. m edges, n vertices, source s , sink t
edge costs c_e , integer in $[-C, C]$, where $C = m^{O(1)}$

Goal: Route 1 unit of
flow from $s \rightarrow t$,
minimize cost

Real-weighted shortest
st-path




Cost = -2

Uncapacitated Min-Cost Flow

Directed graph $G = (V, E)$. m edges, n vertices, source s , sink t
edge costs c_e , integer in $[-C, C]$, where $C = m^{O(1)}$

Goal: Route 1 unit of
flow from $s \rightarrow t$,
minimize cost

Real-weighted shortest
st-path




Cost = -3

Uncapacitated Min-Cost Flow

Directed graph $G = (V, E)$. m edges, n vertices, source s , sink t
edge costs c_e , integer in $[-C, C]$, where $C = m^{O(1)}$

Goal: Route 1 unit of
flow from $s \rightarrow t$,
minimize cost

Real-weighted shortest
st-path



Cost = -4

Linear Algebraic View


- $f \in \mathbb{R}^E$, i.e. a real vector on the edges

Goal: minimize
 $c^T f$

Flow routes the demand:

Net flow on s is +1

Net flow on t is -1



Flow conservation:
all other vertices have
incoming flow=outgoing flow

Linear Program

$$\min_f c^T f = \sum_e c_e f_e$$

Min Cost

For all edges e

$$f_e \geq 0$$

Direction constraints

For all vertices x

$$B^T f = \chi_{s,t}$$

Unit flow from s to t

Transshipment Primal LP

$$\min_f c^T f = \sum_e c_e f_e$$

Min Cost

For all edges e

$$f_e \geq 0$$

Direction constraints

For all vertices x

$$B^T f = d$$

Route the demand

Transshipment Primal LP

$$\min_f c^T f = \sum_e c_e f_e$$

Min Cost

For all edges e

$$f_e \geq 0$$

Direction constraints

For all vertices x


$$B^T f = d$$

Route the demand

Generalize weighted bipartite matching,
capacitated min-cost flow, shortest path, max
flow, ...


Edge insertions: Flow stays feasible

- Add an edge, put 0 flow on it
- Flow stays feasible




Edge insertions: Flow stays feasible

- Add an edge, put 0 flow on it
- Flow stays feasible
- Then move towards smaller cost
- Incremental min-cost flow in almost-linear time
 - [C-Kyng-Liu-Meierhans-Probst Gutenberg '24]
 - Maintain a primal solution under relaxing updates



Edge deletions: Flow no longer feasible

- Not an st -path anymore
- Edge deletion does not relax the primal problem
- How about the dual?



Transshipment Dual LP

$$\max_{y \in \mathbb{R}^V} d^T y$$

For any edge $e = (u, v)$ $c_e \geq y_u - y_v$

Triangle Inequality
Non-negative Slack

Transshipment Dual LP

$$\max_{y \in \mathbb{R}^V} d^T y$$

For any edge $e = (u, v)$ $c - By \geq 0$

Triangle Inequality
Non-negative Slack

- Remove an edge \rightarrow remove one constraint
- Current dual solution y stays feasible
- Move towards larger $d^T y$

Decremental Thresholded Transshipment

Theorem [Brand-C-Kyng-Liu-Meierhans-Probst Gutenberg-Sachdeva]

Given a decremental graph G , threshold D^*

After each edge deletion, either

- Certify that $\max_{c-By \geq 0} d^T y < D^*$, or
- Output a feasible dual potential y s.t. $d^T y \geq D^*$

in $m^{1+o(1)}$ total time deterministically

Decremental Thresholded Transshipment

Theorem [Brand-C-Kyng-Liu-Meierhans-Probst Gutenberg-Sachdeva]

Given a decremental graph G , threshold D^*

After each edge deletion, either

- Certify that $\max_{c-By \geq 0} d^T y < D^*$, or
- Output a feasible dual potential y s.t. $d^T y \geq D^*$

in $m^{1+o(1)}$ total time deterministically

Application: Decremental...

1. max flow
2. min-cost flow
3. st-shortest path
4. single-source reachability
5. SCC maintenance
6. ...

Dual L1 IPM

[Karmarkar '84]

Maximize $d^T y$

Enforce $c - By \geq 0$

$$\min_{y \in R^V} \Phi(y) = 100m \log(D^* - d^T y) + \sum_e -\log(c_e - (y_u - y_v))$$

- If $\Phi(y) \leq -1000m \log m$, $d^T y \geq D^* - m^{-10}$
- Start at $y^{(0)}$ with $\Phi(y^{(0)}) = O(m \log m)$
- Iteratively, $\Phi(y^{(t+1)}) \leq \Phi(y^{(t)}) - m^{-o(1)}$
 - After $T = m^{1+o(1)}$ iterations, $\Phi(y^{(T)}) \leq -1000m \log m$

Dual L1 IPM

[Karmarkar '84]

Maximize $d^T y$

Enforce $c - By \geq 0$

$$\min_{y \in R^V} \Phi(y) = 100m \log(D^* - d^T y) + \sum_e -\log(c_e - (y_u - y_v))$$

- $\Phi(y + \Delta) \approx \Phi(y) + \langle \nabla \Phi, \Delta \rangle + \|UB\Delta\|_1^2$, $U = \text{diag}(c - By)^{-1} = \sqrt{\nabla^2 \Phi}$

- If $\max_{c-By \geq 0} d^T y \geq D^*$, $\min_{\Delta \in R^V} \frac{\langle \nabla \Phi, \Delta \rangle}{\|UB\Delta\|_1} \leq -0.1$

Min-Ratio Cut

- If $\frac{\langle \nabla \Phi, \Delta \rangle}{\|UB\Delta\|_1} \leq -\kappa$, $\Phi(y + \eta\Delta) \leq \Phi(y) - \kappa^2$

$\langle 1, \nabla \Phi \rangle = 0$

Decremental Transshipment Algorithm


- Start at $y^{(0)}$ with $\Phi(y^{(0)}) = O(m \log m)$
- For $t = 0, 1, \dots, T = m^{1+o(1)}$, do
 - $m^{o(1)}$ -approximately minimize $\min_{\Delta \in R^V} \frac{\langle \nabla \Phi, \Delta \rangle}{\|UB\Delta\|_1}$
 - If the ratio is larger than $-1/m^{o(1)}$, **certify** $\max\langle d, y \rangle < D^*$
 - Otherwise, update $y^{(t+1)} = y^{(t)} + \eta\Delta$
- Output $y^{(T)}$

can view it as a L1 trust
region Newton method

Min-Ratio Cut

$$\langle 1, g \rangle = 0$$

$$\min_{\Delta \in R^V} \frac{\langle g, \Delta \rangle}{\|UB\Delta\|_1}$$




$$\langle g, \Delta \rangle = 5 + 0 = 5$$

$$\|UB\Delta\|_1 = 1 + 4 + 3 = 8$$

Min-Ratio Cut

$$\langle 1, g \rangle = 0$$

$$\min_{\Delta \in R^V} \frac{\langle g, \Delta \rangle}{\|UB\Delta\|_1}$$



$$\langle g, \Delta \rangle = -3$$

$$\|UB\Delta\|_1 = 3 + 8 = 11$$


Min-Ratio Cut

$$\langle 1, g \rangle = 0$$

$$\min_{\Delta \in \mathbb{R}^V} \frac{\langle g, \Delta \rangle}{\|UB\Delta\|_1}$$

Optimal solution is a cut, i.e.,

$$\Delta \in \{0,1\}^V$$



$$\langle g, \Delta \rangle = -3$$

$$\|UB\Delta\|_1 = 3 + 8 = 11$$

Fully-Dynamic Min-Ratio Cut

Theorem [Brand-C-Kyng-Liu-Meierhans-Probst Gutenberg-Sachdeva]

Given fully dynamic graph G , there's a **deterministic** data structure that


- $m^{o(1)}$ -approximately minimize $\min_{\Delta \in \mathbb{R}^V} \frac{\langle \nabla \Phi, \Delta \rangle}{\|UB\Delta\|_1}$, and
- Update $y := y + \Delta$, gradients $\nabla \Phi$ and weights $U = \text{diag}(c - By)^{-1}$

In $m^{o(1)}$ time per operation

Approx Min-Ratio Cuts via Tree Cut Sparsifiers


- Approximate the cut structure of G by a tree T
 - [Räcke '08, Mądry '10, Räcke-Shah-Täubig '14, Goranci-Räcke-Saranurak-Tan '21]
- Solve Min-Ratio Cut on T
 - Claim: it cuts one tree edge

Min-ratio cut on a tree cuts one edge




Min-ratio cut on a tree cuts one edge

Min-ratio cut Δ^*




Min-ratio cut on a tree cuts one edge




$$\|UB\Delta^*\|_1 = \sum_{\text{tree cuts } C} \|UBC\|_1$$

Min-ratio cut on a tree cuts one edge




$$\|UB\Delta^*\|_1 = \sum_{\text{tree cuts } C} \|UBC\|_1$$

Min-ratio cut on a tree cuts one edge




$$\|UB\Delta^*\|_1 = \sum_{\text{tree cuts } C} \|UBC\|_1$$

Min-ratio cut on a tree cuts one edge



$$\|UB\Delta^*\|_1 = \sum_{\text{tree cuts } C} \|UBC\|_1$$

One of  is as good as Δ^*


G is not always a tree

- G is a (weighted) ϕ -expander if

$$\frac{U(S, V \setminus S)}{\text{vol}_G(S)} \geq \phi, \forall S \subseteq V$$

- One of the singleton cut is $(1/\phi)$ -approx.


$$\phi \sum_{u \in S} \deg(u) \leq U(S, V \setminus S) \leq \sum_{u \in S} \deg(u)$$



Tree Cut Sparsifier via Expander Hierarchy

[Goranci-Räcke-Saranurak-Tan '21]


- Weighted Expander Decomposition
 - Decompose G into $m^{o(1)}$ -expanders $G[V_1], G[V_2], \dots$ such that
Inter-expander edge total weight $\leq \phi \cdot m^{o(1)} \cdot \text{total edge weight}$



Tree Cut Sparsifier via Expander Hierarchy

[Goranci-Räcke-Saranurak-Tan '21]


- Weighted Expander Decomposition
- Approx each expander by a star



Tree Cut Sparsifier via Expander Hierarchy

[Goranci-Räcke-Saranurak-Tan '21]


- Weighted Expander Decomposition
- Approx each expander by a star
- Contract



Tree Cut Sparsifier via Expander Hierarchy

[Goranci-Räcke-Saranurak-Tan '21]


- Weighted Expander Decomposition
- Approx each expander by a star
- Contract
- Repeat until one vertex left



Tree Cut Sparsifier via Expander Hierarchy

[Goranci-Räcke-Saranurak-Tan '21]

- Weighted Expander Decomposition
- Approx each expander by a star
- Contract
- Repeat until one vertex left
- Output T = union of stars
 - $m^{o(1)}$ -approx



Issue


- How to Weighted Expander Decomposition?

- Approx each expander by a star

- Contract

- Repeat until one vertex left

- Output T = union of stars



Weighted Expander Decomposition

- Previous approach is hard to dynamize [\[Li-Saranurak '21\]](#)
- Reduce it to unweighted case via rounding

Weighted Expander Decomposition

- Initialize $G' = (V, \emptyset)$
- Let U^{total} be the total edge weight.
- For any edge e with $u(e) \geq \frac{\phi U^{total}}{m}$, add $\left\lceil \frac{u(e) \cdot m}{\phi U^{total}} \right\rceil$ parallel edges to G'
- For each vertex u , add $\left\lceil \frac{\deg(u) \cdot m}{\phi U^{total}} \right\rceil$ self-loops to G'
- Run unweighted expander decomposition on G' and output the partition

Weighted Expander Decomposition

Can be made dynamic

- Initialize $G' = (V, \emptyset)$
- Let U^{total} be the total edge weight.
- For any edge e with $u(e) \geq \frac{\phi U^{total}}{m}$, add $\left\lceil \frac{u(e) \cdot m}{\phi U^{total}} \right\rceil$ parallel edges to G'
- For each vertex u , add $\left\lceil \frac{\deg(u) \cdot m}{\phi U^{total}} \right\rceil$ self-loops to G'
- Run unweighted expander decomposition on G' and output the partition

Fully-Dynamic Weighted Expander Hierarchy

Theorem [Brand-C-Kyng-Liu-Meierhans-Probst Gutenberg-Sachdeva]

Given fully dynamic weighted graph G ,
there is a deterministic algorithm that
maintains an expander hierarchy with
 $m^{o(1)}$ update time.

Fully-dynamic $m^{o(1)}$ -approx

1. All-pair maxflow/mincut
2. Sparsest cuts
3. K-comm flows
4. Vertex cut sparsifier
5. ...

Summary

- Decremental Min-cost flow via maintaining dual solution
- Use L1 IPM to gradually improve/certify the current solution
- Per step, solve a min-ratio cut problem
- Weighted expander hierarchy: tree approximation of graphs w.r.t. cuts

Future direction

- $T_{maxflow}/m$: $\exp(\log^{7/8}m) \rightarrow \exp(\log^{3/4}m) \rightarrow \text{polylog?}$
 - Deterministic: $\exp(\log^{17/18}m) \rightarrow \exp(\log^{5/6}m) \rightarrow ??$
- Decremental $(1 + \varepsilon)$ -approx. directed SSSP?
- Implication on general graph matching?

- $T_{maxflow}/m$: $\exp(\log^{7/8}m) \rightarrow \exp(\log^{3/4}m) \rightarrow \text{polylog?}$

- Deterministic: $\exp(\log^{17/18}m) \rightarrow \exp(\log^{5/6}m) \rightarrow ??$

Thanks!!

- Decremental $(1 + \varepsilon)$ -approx. directed SSSP?

- Implication on general graph matching?



Jan van den
Brand
Georgia Tech



Rasmus Kyng
ETH



Yang P. Liu
IAS



Simon
Meierhans
ETH



Maximilian
Probst Gutenberg
ETH



Sushant
Sachdeva
U. Toronto